

Deva Sri Sasikala

Software Engineer

Virginia | devasrisasikala3@gmail.com | (571) 210-0132 | [LinkedIn](#)

PROFESSIONAL SUMMARY

Software Engineer with **5+ years** of experience designing and building **scalable distributed systems, microservices, and cloud-native applications** across **fintech, e-commerce, and enterprise platforms**. Expertise in **Java, Python, Spring Boot, Apache Kafka, Apache Spark, AWS, Kubernetes, SQL, and Distributed Systems**, delivering high-performance backend services handling **10K+ RPS** and **millions of daily transactions** with **sub-200ms latency**. Proven track record improving **system reliability, scalability, and developer productivity** through **event-driven architectures, CI/CD automation, observability, LLM, RAG, and AI-powered platform integrations** while maintaining **99.9%+ availability**.

TECHNICAL SKILLS

Languages: Java, Python, PySpark

Backend & APIs: Spring Boot, Spring Batch, Hibernate, REST APIs, gRPC, Microservices, Event-Driven Architecture, Distributed Systems

AI/ML & MLOps: Mosaic AI, RAG Pipelines, LLM Services, OpenAI API, Prompt Engineering, NLP, Semantic Search, Embeddings, Vector Search, Information Retrieval, MLflow, Model Serving, Feature Engineering, Anomaly Detection

Big Data & Pipelines: Apache Spark, Databricks, Delta Lake, ETL Orchestration, Spring Batch ETL

Databases & Caching: PostgreSQL, Redis, Caffeine, Vector Databases

Messaging & Streaming: Apache Kafka

Cloud & DevOps Infrastructure: AWS (EC2, RDS, S3, IAM, KMS, Lambda), Azure Blob Storage, Docker, Kubernetes, Jenkins, CI/CD Pipeline

Observability: ELK Stack, Prometheus, Grafana, Centralized Logging

Security & Compliance: KYC, Identity Verification, AML Screening, IAM, KMS

Data Structures & Algorithms: LRU Cache, Hash Maps, Priority Queues

EXPERIENCE

Databricks | CA | Software Engineer

January 2025 – Present

- Built an end-to-end distributed **AI-powered** enterprise knowledge platform using **Mosaic AI, RAG pipelines, vector search, and LLM** services, supporting millions of daily queries with **sub-200ms p95 latency** and **99.9%+ availability**.
- Designed and implemented **high-throughput backend microservices in Python**, supporting **10K+ RPS peak load**, with resilient API design, distributed system patterns, and fault-tolerant LLM inference integration.
- Engineered **ML and MLOps pipelines on Databricks** using **Apache Spark, PySpark, MLflow**, and **CI/CD automation**, reducing **deployment time by 60%** and accelerating **train-to-prod iteration cycles by 3 times** across **100+** experiments.
- Developed a **semantic search and retrieval system** leveraging **embeddings, vector databases, and IR techniques**, improving **search relevance (CTR +25–35%)** and reducing **query latency by 40%**.
- Built scalable **cloud-native data pipelines** using **Delta Lake, AWS S3, Azure Blob Storage**, and **ETL orchestration**, enabling **petabyte-scale processing**, reducing pipeline runtime by **30%**, and improving data freshness from daily to near real-time (hourly).
- Implemented **event-driven, observable backend systems** using **Kafka, gRPC, Kubernetes (EKS)** and centralized logging, improving system scalability by **25%**, reducing **MTTD by 30–40%**, and achieving near-zero unplanned downtime.

SoFi | CA | Software Engineer

March 2023 – December 2024

- Engineered and deployed a real-time **Fraud Detection Platform** using **Python, Spring Boot, Machine Learning, gRPC**, and **REST APIs** with model serving and risk scoring, enabling **sub-200ms transaction** monitoring and cutting fraud losses by **22%**.
- Built streaming data pipelines with **Apache Kafka, Apache Spark, and Redis**, applying **feature engineering**, event processing, and behavioral analytics to power anomaly detection and reduce false-positive alerts by **18%**.
- Developed **LLM-powered risk investigation systems** using **OpenAI API, prompt engineering, NLP, semantic search, and knowledge retrieval**, cutting manual case-review effort by **30%**.
- Scaled high-throughput payment processing services on a **microservices** and **event-driven architecture** with **Spring Boot, Kafka, and PostgreSQL**, reducing settlement latency by **35% at 99.8% uptime**.
- Owned **cloud-native infrastructure** and **SRE** across **AWS (IAM, KMS, Lambda), Docker, Kubernetes**, and **CI/CD**, building observability (**metrics, logging, alerting**) that accelerated incident resolution and eliminated recurring security vulnerabilities.

Myntra | India | Software Engineer

January 2020 – November 2021

- Built an end-to-end E-commerce order management platform using **Java, Spring Boot, Apache Kafka, AWS (EC2, RDS, S3)**, and **REST APIs**, supporting catalog, inventory, and order workflows across **500K+ SKUs** and **1M+ daily transactions** with **99.3% order success rate**.
- Designed high-throughput **REST** and **gRPC** services backed by **PostgreSQL, Hibernate, and Redis**, implementing **Caffeine caching** to reduce database load by **25%**, sustain **60K+ concurrent users**, and improve API latency by **20%**.
- Engineered **Spring Batch** reconciliation pipelines processing **400K+ daily records** with **99.9% data accuracy**, while optimizing payment processing through efficient **data structures and algorithms (LRU Cache, Hash Maps, Priority Queues)** to support **2x traffic spikes**.
- Automated **CI/CD** pipelines using **Jenkins, AWS CloudFormation, and Docker**, delivering **20+ zero-downtime deployments** per week and reducing **MTTR by 35%** through monitoring and observability with **ELK Stack, Prometheus, and Grafana**.

EDUCATION

Ph.D. in Information Technology (AI Concentration) | University of the Cumberland, KY (*Expected 2029*)

June 2026 – Present

Master's in Computer Science | University of Missouri-Kansas City, Missouri, USA.

January 2022 – May 2023